

Network setting

K3s will come with pretty much everything pre-configured like `traefik`.

More info about `traefik`: <https://doc.traefik.io/traefik/>

However, I would like to have LoadBalancer, and in essence to be able to give services (pods) an external IP, just like my Kubernetes nodes, not from internal Kubernetes ranges. Normally, this is an external component, and your cloud provider should somehow magically give that to you, but since we are our own cloud provider, and we are trying to keep everything in one cluster... In short MetalLB is the answer.

What is MetalLB

<https://metalb.universe.tf/>

Deployment

This is a two-step process: we deploy MetalLB load balancer, and then push configuration to it, and tell it what range of IPs to use.

Apply the following: first we will create a namespace called `metalb-system`, and second we will deploy MetalLB into it.

Note

Look here <https://metalb.universe.tf/installation/> for the most up to date version of metalb links.

kubectl apply -f https://raw.githubusercontent.com/metalb/metalb/v0.9.5/manifests/namespace.yaml
kubectl apply -f https://raw.githubusercontent.com/metalb/metalb/v0.9.5/manifests/metalb.yaml

Note

I prefer to store configuration files in folders named after components we deploy to cluster so that I can easily delete the service later without looking at links pointing to the Internet. So, you can just create folder `MetalLB`, and download the yaml file into it for later use.

We need to create a secret key for the `speakers` (the MetalLB pods) to encrypt speaker communications:

```
kubectl create secret generic -n metalb-system memberlist --from-literal=secretkey="$(openssl rand -base64 128)"
```

Configuration

Next, create `config.yaml` in your `MetalLB` folder, here we are going to tell MetalLB what IPs to use:

```
apiVersion: v1
kind: ConfigMap
metadata:
  namespace: metalb-system
  name: config
data:
  config: |
    address-pools:
    - name: default
      protocol: layer2
      addresses:
      - 192.168.0.238-192.168.0.250
```

As you can see, I specified a range from 192.168.0.230 to 192.168.0.250. That will give me 20 "external" IPs to work with for now.

Apply the config:

```
kubectl apply -f config.yaml
```

Check

Check if everything deployed OK

```
root@control01:~# kubectl get pods -n metalb-system
NAME                READY   STATUS    RESTARTS   AGE
controller-65db86dc6-7h59v    1/1     Running   0           6d5h
speaker-6vjzn             1/1     Running   0           6d5h
speaker-b25rk             1/1     Running   0           6d5h
speaker-dw2pv             1/1     Running   0           6d5h
speaker-gdjr              1/1     Running   0           6d5h
speaker-hc72j             1/1     Running   0           6d5h
speaker-k9nzq             1/1     Running   0           6d5h
speaker-mfmkq             1/1     Running   0           6d5h
speaker-qvzv2             1/1     Running   0           6d5h
speaker-z6dk6             1/1     Running   0           6d5h
```

You should have as many `speaker-xxxx` as you have nodes in the cluster, since they run one per node.

Now services that use LoadBalancer should have an external IP assigned to them.

For example:

```
root@control01:~# kubectl get svc --all-namespaces
NAMESPACE   NAME          TYPE          CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE
default     kubernetes    ClusterIP     10.43.0.1    <none>        443/TCP          8d
kube-system kube-dns      ClusterIP     10.43.0.18   <none>        53/UDP, 53/TCP, 9153/TCP  8d
kube-system metrics-server ClusterIP     10.43.246.167 <none>        443/TCP          8d
kube-system traefik      LoadBalancer  10.43.61.64    192.168.0.238 80:31712/TCP, 443:31124/TCP  8d
kube-system traefik-prometheus ClusterIP     10.43.178.172 <none>        9180/TCP          8d
kubernetes-dashboard dashboard-metrics-scraper ClusterIP     10.43.147.66   <none>        8080/TCP          8d
kubernetes-dashboard kubernetes-dashboard ClusterIP     10.43.151.258 <none>        443/TCP          8d
longhorn-system csi-attacher  ClusterIP     10.43.125.73   <none>        12345/TCP         8d
longhorn-system csi-provisioner ClusterIP     10.43.118.73   <none>        12345/TCP         8d
longhorn-system csi-resizer   ClusterIP     10.43.245.224 <none>        12345/TCP         8d
longhorn-system csi-snapshotter ClusterIP     10.43.238.3    <none>        12345/TCP         8d
longhorn-system longhorn-backend ClusterIP     10.43.118.82   <none>        9500/TCP          8d
longhorn-system longhorn-frontend ClusterIP     10.43.284.227 <none>        80/TCP           8d
openfaas alertmanager ClusterIP     10.43.79.38    <none>        9093/TCP          6d6h
openfaas basic-auth-plugin ClusterIP     10.43.163.133 <none>        8080/TCP          6d6h
openfaas gateway ClusterIP     10.43.187.155 <none>        8080/TCP          6d6h
openfaas gateway-external NodePort      10.43.71.53    <none>        8080:31112/TCP    6d6h
openfaas nats ClusterIP     10.43.122.99   <none>        4222/TCP          6d6h
openfaas prometheus ClusterIP     10.43.78.247   <none>        9090/TCP          6d6h
```

Look how `traefik` automatically got an IP from the external range. In the end, this is what you would want: Not to point to a single node IP and be redirected based on DNS, which would stop working the moment the node with that IP died. This way, we make the external IP node independent. Now, you can point DNS to that IP and be sure it will be routed correctly.

Note

This is how I prefer my network settings, and makes most sense to me when creating external services. I'm sure there are like a hundred different methods using external load balancers, Nginx ingress (basically reverse proxy) and who knows what in production, but hey, there is no official one standardized setting for Kubernetes (which can be such a pain sometimes) so who's to say this is not OK? 😊

🍷 Liked it ? Buy me a drink :)

Last update: October 20, 2021

Comments

What do you think?

6 Responses

👍

🤔

😍

😮

😡

😞

Upvote

Funny

Love


Surprised


Angry

Sad

7 Comments

<https://rpi4cluster.com>

 Disqus' Privacy Policy


 Login

♥ Favorite

🐦 Tweet

📌 Share

Sort by Best

 Join the discussion...

LOG IN WITH

OR SIGN UP WITH DISQUS ?

Name

Alex Ellis • 6 months ago

Have you tried the inlets-operator yet? It works really well with MetalLB and Traefik to give you public LoadBalancer IPs alongside any local ones you want to keep.

No messy NAT or port-forwarding rules required. <https://blog.zespre.com/inl...>

1 ^ | v • Reply • Share ›

vladoportos Mod ➔ Alex Ellis • 6 months ago

Not yet, I knew about it from you, I think you have introduced the concept to me.. Either on blog or in your e-book (can't remember now). They are definitely a good way to expose service, but I did not want that at the time. I'm in the process of rebuilding my cluster from scratch right now. Including design of new enclosure for RPIs and Power system.. So as soon as the enclosure is done, I will be reinstalling everything again and improving this guide as well to be more up to date. I will add inlets as well.

1 ^ | v • Reply • Share ›

Freman • 2 months ago

So, I've done this 6 or 7 times now and it worked out perfectly, but this last time I moved my kubes cluster onto it's own vlan at last and started rebuilding, only I forgot to change my config from 10.0.0.10-10.0.0.30 to 10.0.1.200-10.0.1.250... now my traefik is on 10.0.0.30 and I can't figure out how to make it move (even after I changed configs)

^ | v • Reply • Share ›

Karthik • a year ago

I had a nightmare with metalb and k3s. It seems k3s has inbuilt service that assigns IP to loadbalancer type services.

Check this

<https://github.com/metalb/...>

Might be useful to someone.

btw, your guide is really detailed and very much helpful for those who start with k3s like me. thanks.

^ | v • Reply • Share ›

vladoportos Mod ➔ Karthik • a year ago

Yes the issue is known, I did use --disable servicelb when installing K3s :)

^ | v • Reply • Share ›

Devin Flake ➔ vladoportos • 6 months ago

Do you know if this is still an issue? I've been playing around with k3s and MetalLB but I'm not experienced enough to find any problems.

^ | v • Reply • Share ›