

Install Service Monitors

Longhorn Servicemonitor

Our storage provisioner Longhorn, that we deployed somewhere near the start of this whole K3s Kubernetes cluster setup, also natively provides data for Prometheus.

Create a new folder, `monitoring`, that we will put most of our configs in, and create the file `longhorn-servicemonitor.yaml`.

```
cd
mkdir monitoring
cd monitoring
touch longhorn-servicemonitor.yaml
```

Edit the `longhorn-servicemonitor.yaml`:

```
apiVersion: monitoring.coreos.com/v1
kind: ServiceMonitor
metadata:
  name: longhorn-prometheus-servicemonitor
  namespace: monitoring
  labels:
    name: longhorn-prometheus-servicemonitor
spec:
  selector:
    matchLabels:
      app: longhorn-manager
    namespaceSelector:
      matchNames:
        - longhorn-system
  endpoints:
    - port: manager
```

As you can see, we are not talking to Kubernetes API (we are.. but..), but to `apiVersion: monitoring.coreos.com/v1`, so we are basically telling Prometheus Operator to create something for us. In this case it's kind: `ServiceMonitor`.

Makes sense, right? Next, `metadata`, and here I'm not 100% sure about the `name`: and below `labels`: -> `name`:. I know that we refer to these later, or one of them, when we tell Prometheus which Service Monitor to collect data from.

This should be clear, `metadata`: -> `namespace: monitoring`, we are telling it to deploy into our monitoring namespace.

The rest under `spec`: is basically telling what app the Service Monitor should "bind to". It's looking for `app: longhorn-manager` in namespace `longhorn-system` and `port: manager`. This port could be a port number, but it also can have a name, so in this case it's named `manager`.

This is the `longhorn-manager` we are targeting.

```
root@control01:/home/ubuntu/monitoring# kubectl get daemonset -n longhorn-system
NAME                                DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE   NODE SELECTOR   AGE
-
-
longhorn-manager                    9         9         9       9            9           <none>         35d
```

And if you try to describe it with:

```
kubectl describe daemonset longhorn-manager -n longhorn-system
-
-
-
Port:          9500/TCP
-
-
-
```

You get that the port it is using is 9500/TCP, but I don't know where it's set that `manager == 9500`. If you know, please comment below.

Node-exporter

This is the daemon set we will deploy to collect metrics from individual cluster nodes, underlying HW, etc...

In the monitoring folder, create a new folder called 'node-exporter' and create the following files in it:

1

Info

Sorry guys, I always try to provide a source for the code, but for the life of me I can't find where I got this deployment from. If you have seen it before somewhere, let me know down in the comments, and I will update the source here.

cluster-role-binding.yaml

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: node-exporter
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: node-exporter
subjects:
- kind: ServiceAccount
  name: node-exporter
  namespace: monitoring
```

cluster-role.yaml

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: node-exporter
rules:
- apiGroups:
  - authentication.k8s.io
  resources:
  - tokenreviews
  verbs:
  - create
- apiGroups:
  - authorization.k8s.io
  resources:
  - subjectaccessreviews
  verbs:
  - create
```

service-account.yaml

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: node-exporter
  namespace: monitoring
```

service.yaml

```
apiVersion: v1
kind: Service
metadata:
  labels:
    app.kubernetes.io/name: node-exporter
  name: node-exporter
  namespace: monitoring
spec:
  clusterIP: None
  ports:
  - name: https
    port: 9100
    targetPort: https
  selector:
    app.kubernetes.io/name: node-exporter
```

daemonset.yaml

```
apiVersion: apps/v1
kind: DaemonSet
metadata:
  labels:
    app.kubernetes.io/component: exporter
    app.kubernetes.io/name: node-exporter
    app.kubernetes.io/part-of: kube-prometheus
    app.kubernetes.io/version: 1.1.0
  name: node-exporter
  namespace: monitoring
spec:
  selector:
    matchLabels:
      app.kubernetes.io/component: exporter
      app.kubernetes.io/name: node-exporter
      app.kubernetes.io/part-of: kube-prometheus
  template:
    metadata:
      labels:
        app.kubernetes.io/component: exporter
        app.kubernetes.io/name: node-exporter
        app.kubernetes.io/part-of: kube-prometheus
        app.kubernetes.io/version: 1.1.0
    spec:
      containers:
```

```
- args:
  --web-listen-address=127.0.0.1:9100
  --path.sysfs=/host/sys
  --path.rootfs=/host/root
  --no-collector.wifi
  --no-collector.hamon
  --collector.filesystem.ignored-mount-points='^(dev|proc|sys|var/lib/docker/.+|var/lib/kubelet/pods/.+)($)'
  --collector.netclass.ignored-devices=(veth.*)$
  --collector.netdev.device-exclude=(veth.*)$
image: quay.io/prometheus/node-exporter:v1.1.0
name: node-exporter
resources:
  limits:
    cpu: 250m
    memory: 180Mi
  requests:
    cpu: 102m
    memory: 180Mi
volumeMounts:
- mountPath: /host/sys
  mountPropagation: HostToContainer
  name: sys
  readOnly: true
- mountPath: /host/root
  mountPropagation: HostToContainer
  name: root
  readOnly: true
- args:
  --logtostderr
  --secure-listen-address=[$(IP)]:9100
  --tls-cipher-
suite=TLS_ECDHE_RSA_WITH_AES_128_GCM_SHA256,TLS_ECDHE_ECDSA_WITH_AES_128_GCM_SHA256,TLS_ECDHE_RSA_WITH_AES_256_GCM_SHA384,TLS_ECDHE_ECDSA_WITH_AES_256_GCM_SHA384,TLS_ECDH_ECDSA_WITH_AES_256_GCM_SHA384,TLS_ECDH_RSA_WITH_AES_256_GCM_SHA384,TLS_DHE_RSA_WITH_AES_256_GCM_SHA384,TLS_DHE_DSS_WITH_AES_256_GCM_SHA384,TLS_ECDHE_RSA_WITH_
  --upstream=http://127.0.0.1:9100/
env:
- name: IP
  valueFrom:
    fieldRef:
      fieldPath: status.podIP
image: quay.io/brancz/kube-rbac-proxy:v0.8.0
name: kube-rbac-proxy
ports:
- containerPort: 9100
  hostPort: 9100
  name: https
resources:
  limits:
    cpu: 20m
    memory: 40Mi
  requests:
    cpu: 10m
    memory: 20Mi
securityContext:
  runAsGroup: 65532
  runAsNonRoot: true
  runAsUser: 65532
hostNetwork: true
hostPID: true
nodeSelector:
  kubernetes.io/os: linux
securityContext:
  runAsNonRoot: true
  runAsUser: 65534
serviceAccountName: node-exporter
tolerations:
- operator: Exists
volumes:
- hostPath:
    path: /sys
    name: sys
- hostPath:
    path: /
    name: root
updateStrategy:
  rollingUpdate:
    maxUnavailable: 10%
  type: RollingUpdate
```

service-monitor.yaml

```
apiVersion: monitoring.coreos.com/v1
kind: ServiceMonitor
metadata:
  labels:
    app.kubernetes.io/name: node-exporter
  name: node-exporter
  namespace: node-exporter
spec:
  endpoints:
    - bearerTokenFile: /var/run/secrets/kubernetes.io/serviceaccount/token
      interval: 15s
      port: https
      relabelings:
        - action: replace
          regex: (.*)
          replacement: $1
          sourceLabels:
            - __meta_kubernetes_pod_node_name
          targetLabel: instance
      scheme: https
      tlsConfig:
        insecureSkipVerify: true
      jobLabel: app.kubernetes.io/name
    selector:
      matchLabels:
        app.kubernetes.io/name: node-exporter
```

You can deploy all the yaml files by going in folder above and apply -f on the folder.

```
cd ..
kubectl apply -f node-exporter/
```

This will create all permissions, and deploy the pod with the application Node Exporter, that will read metrics from Linux.

After doing so, you should see `node-exporter-xxxx` pods in the `monitoring` namespace; I have 9 nodes so it's there 9 times.

```
root@control01: /home/ubuntu/monitoring/node-exporter# kubectl get pods -n monitoring
```

NAME	READY	STATUS	RESTARTS	AGE
node-exporter-8ttd6	2/2	Running	0	15d
node-exporter-cfghh	2/2	Running	0	15d
node-exporter-gdr7f	2/2	Running	0	15d
node-exporter-kkd68	2/2	Running	0	15d
node-exporter-mmk45	2/2	Running	0	15d
node-exporter-np6q5	2/2	Running	0	15d
node-exporter-pfmcn	2/2	Running	2	15d
node-exporter-rdlwn	2/2	Running	0	15d
node-exporter-rvz61	2/2	Running	0	15d

Kube State Metrics

This is a simple service that listens to the Kubernetes API, and generates metrics about the state of the objects.

Link to official github: [kube-state-metrics](#) .

Again as before, create a new folder in our monitoring folder, called `kube-state-metrics`. Create the following files in it

kube-state-metrics-clusterRole.yaml

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  labels:
    app.kubernetes.io/component: exporter
    app.kubernetes.io/name: kube-state-metrics
    app.kubernetes.io/part-of: kube-prometheus
    app.kubernetes.io/version: 1.9.7
  name: kube-state-metrics
rules:
- apiGroups:
  - ""
  resources:
  - configmaps
  - secrets
  - nodes
  - pods
  - services
  - resourcequotas
  - replicationcontrollers
  - limitranges
  - persistentvolumeclaims
  - persistentvolumes
  - namespaces
  - endpoints
verbs:
- list
- watch
- apiGroups:
  - extensions
  resources:
  - daemonsets
  - deployments
  - replicasets
  - ingresses
verbs:
- list
- watch
- apiGroups:
  - apps
```

```
resources:
- statefulsets
- daemonsets
- deployments
- replicaset
verbs:
- list
- watch
- apiGroups:
- batch
resources:
- cronjobs
- jobs
verbs:
- list
- watch
- apiGroups:
- autoscaling
resources:
- horizontalpodautoscalers
verbs:
- list
- watch
- apiGroups:
- authentication.k8s.io
resources:
- tokenreviews
verbs:
- create
- apiGroups:
- authorization.k8s.io
resources:
- subjectaccessreviews
verbs:
- create
- apiGroups:
- policy
resources:
- poddisruptionbudgets
verbs:
- list
- watch
- apiGroups:
- certificates.k8s.io
resources:
- certificatesigningrequests
verbs:
- list
- watch
- apiGroups:
- storage.k8s.io
resources:
- storageclasses
- volumeattachments
verbs:
- list
- watch
- apiGroups:
- admissionregistration.k8s.io
resources:
- mutatingwebhookconfigurations
- validatingwebhookconfigurations
verbs:
- list
- watch
- apiGroups:
- networking.k8s.io
resources:
- networkpolicies
verbs:
- list
- watch
```

kube-state-metrics-clusterRoleBinding.yaml

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  labels:
    app.kubernetes.io/component: exporter
    app.kubernetes.io/name: kube-state-metrics
    app.kubernetes.io/part-of: kube-prometheus
    app.kubernetes.io/version: 1.9.7
  name: kube-state-metrics
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: kube-state-metrics
subjects:
- kind: ServiceAccount
  name: kube-state-metrics
  namespace: monitoring
```

kube-state-metrics-serviceAccount.yaml

```
apiVersion: v1
kind: ServiceAccount
metadata:
  labels:
    app.kubernetes.io/component: exporter
    app.kubernetes.io/name: kube-state-metrics
    app.kubernetes.io/part-of: kube-prometheus
    app.kubernetes.io/version: 1.9.7
  name: kube-state-metrics
  namespace: monitoring
```

kube-state-metrics-service.yaml

```
apiVersion: v1
kind: Service
metadata:
  labels:
    app.kubernetes.io/component: exporter
    app.kubernetes.io/name: kube-state-metrics
    app.kubernetes.io/part-of: kube-prometheus
    app.kubernetes.io/version: 1.9.7
  name: kube-state-metrics
  namespace: monitoring
spec:
  clusterIP: None
  ports:
  - name: https-main
    port: 8443
    targetPort: https-main
  - name: https-self
    port: 9443
    targetPort: https-self
  selector:
    app.kubernetes.io/component: exporter
    app.kubernetes.io/name: kube-state-metrics
    app.kubernetes.io/part-of: kube-prometheus
```

kube-state-metrics-deployment.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app.kubernetes.io/component: exporter
    app.kubernetes.io/name: kube-state-metrics
    app.kubernetes.io/part-of: kube-prometheus
    app.kubernetes.io/version: 1.9.7
  name: kube-state-metrics
  namespace: monitoring
spec:
  replicas: 1
  selector:
    matchLabels:
      app.kubernetes.io/component: exporter
      app.kubernetes.io/name: kube-state-metrics
      app.kubernetes.io/part-of: kube-prometheus
  template:
    metadata:
      labels:
        app.kubernetes.io/component: exporter
        app.kubernetes.io/name: kube-state-metrics
        app.kubernetes.io/part-of: kube-prometheus
        app.kubernetes.io/version: 1.9.7
    spec:
      containers:
      - args:
        - --host=127.0.0.1
        - --port=8081
        - --telemetry-host=127.0.0.1
        - --telemetry-port=8082
        image: carlosdp/kube-state-metrics:v1.9.6
        name: kube-state-metrics
        resources:
          limits:
            cpu: 100m
            memory: 250Mi
          requests:
            cpu: 10m
            memory: 100Mi
        - args:
        - --logtostderr
        - --secure-listen-address=:8443
        - --tls-cipher-
          suites=TLS_ECDHE_RSA_WITH_AES_128_GCM_SHA256,TLS_ECDHE_ECDSA_WITH_AES_128_GCM_SHA256,TLS_ECDHE_RSA_WITH_AES_256_GCM_SHA384,TLS_ECDHE_ECDSA_WITH_AES_256_GCM_SHA384,TLS_ECD
          - --upstream=http://127.0.0.1:8881/
        image: quay.io/brancz/kube-rbac-proxy:v0.8.0
        name: kube-rbac-proxy-main
        ports:
```

```
- containerPort: 8443
  name: https-main
  resources:
    limits:
      cpu: 20m
      memory: 40Mi
    requests:
      cpu: 10m
      memory: 20Mi
  securityContext:
    runAsGroup: 65532
    runAsNonRoot: true
    runAsUser: 65532
- args:
  - --logtostderr
  - --secure-listen-address=:9443
  - --tls-cipher=
suites=TLS_ECDHE_RSA_WITH_AES_128_GCM_SHA256,TLS_ECDHE_ECDSA_WITH_AES_128_GCM_SHA256,TLS_ECDHE_RSA_WITH_AES_256_GCM_SHA384,TLS_ECDHE_ECDSA_WITH_AES_256_GCM_SHA384,TLS_ECDHE_RSA_WITH_AES_128_GCM_SHA256,TLS_ECDHE_ECDSA_WITH_AES_128_GCM_SHA256
- --upstream=http://127.0.0.1:8082/
image: quay.io/brancz/kube-rbac-proxy:v0.8.0
name: kube-rbac-proxy-self
ports:
- containerPort: 9443
  name: https-self
  resources:
    limits:
      cpu: 20m
      memory: 40Mi
    requests:
      cpu: 10m
      memory: 20Mi
  securityContext:
    runAsGroup: 65532
    runAsNonRoot: true
    runAsUser: 65532
nodeSelector:
  node-type: worker
serviceAccountName: kube-state-metrics
```

Info

Above you can see I'm using `image: carlosedp/kube-state-metrics:v1.9.6`. This is not the official image. (This is the official [Docker Hub Link](#)), but official image is not compiled for arm64, therefore I have changed it to [Docker Hub Link](#). This is the same software, just compiled for more architectures.

kube-state-metrics-serviceMonitor.yaml

```
apiVersion: monitoring.coreos.com/v1
kind: ServiceMonitor
metadata:
  labels:
    app.kubernetes.io/component: exporter
    app.kubernetes.io/name: kube-state-metrics
    app.kubernetes.io/part-of: kube-prometheus
    app.kubernetes.io/version: 1.9.7
  name: kube-state-metrics
  namespace: monitoring
spec:
  endpoints:
    - bearerTokenFile: /var/run/secrets/kubernetes.io/serviceaccount/token
      honorLabels: true
      interval: 30s
      port: https-main
      relabelings:
        - action: labeldrop
          regex: (pod|service|endpoint|namespace)
      scheme: https
      scrapeTimeout: 30s
      tlsConfig:
        insecureSkipVerify: true
    - bearerTokenFile: /var/run/secrets/kubernetes.io/serviceaccount/token
      interval: 30s
      port: https-self
      scheme: https
      tlsConfig:
        insecureSkipVerify: true
      jobLabel: app.kubernetes.io/name
      selector:
        matchLabels:
          app.kubernetes.io/component: exporter
          app.kubernetes.io/name: kube-state-metrics
          app.kubernetes.io/part-of: kube-prometheus
```

And again, jump one folder up and apply everything:

```
cd ..
kubectl apply -f kube-state-metrics/
```

Check the pods in the `monitoring` namespace if you have `kube-state-metrics-xxx` up and running:

```
root@control01:~/home/ubuntu/monitoring/kube-state-metrics# kubectl get pods -n monitoring
NAME                                READY   STATUS    RESTARTS   AGE
-
-
kube-state-metrics-6f8f5578-p99td   3/3     Running   0           14d
-
-
```

We have two more Service Monitors to go. 🐼

Kubelet

Kubelet, in case you did not know, is an essential part of Kubernetes' control plane, and is also something that exposes Prometheus metrics by default in the port 10255. So, it makes sense to create a Service Monitor for it as well.

'But', you surely ask me, 'I just deployed kube-state-metrics, why the fuck do I need another Kubelet thingy monitor?' Well, kube-state-metrics collects lots of data, and some of them overlap with Kubelet provided metrics, but not all; some information can be collected only from Kubelet.

We only need to create one file: `kubelet-servicemonitor.yaml`:

```
apiVersion: monitoring.coreos.com/v1
kind: ServiceMonitor
metadata:
  labels:
    app.kubernetes.io/name: kubelet
  name: kubelet
  namespace: monitoring
spec:
  endpoints:
    - bearerTokenFile: /var/run/secrets/kubernetes.io/serviceaccount/token
      honorLabels: true
      interval: 30s
      metricRelabelings:
        - action: drop
          regex:
kubelet_(pod_worker_latency_microseconds|pod_start_latency_microseconds|cgroup_manager_latency_microseconds|pod_worker_start_latency_microseconds|plug_relist_latency_microseconds|_total)
      sourceLabels:
        - __name__
      - action: drop
        regex:
scheduler_(e2e_scheduling_latency_microseconds|scheduling_algorithm_predicate_evaluation|scheduling_algorithm_priority_evaluation|scheduling_algorithm_preemption_evaluation|_total)
      sourceLabels:
        - __name__
      - action: drop
        regex:
apiserver_(request_count|request_latencies|request_latencies_summary|dropped_requests|storage_data_key_generation_latencies_microseconds|storage_transformation_failures|_total)
      sourceLabels:
        - __name__
      - action: drop
        regex:
kubelet_docker_(operations|operations_latency_microseconds|operations_errors|operations_timeout)
      sourceLabels:
        - __name__
      - action: drop
        regex:
reflector_(items_per_list|items_per_watch|list_duration_seconds|lists_total|short_watches_total|watch_duration_seconds|watches_total)
      sourceLabels:
        - __name__
      - action: drop
        regex:
etcd_(helper_cache_hit_count|helper_cache_miss_count|helper_cache_entry_count|request_cache_get_latencies_summary|request_cache_add_latencies_summary|request_latencies_summary|_total)
      sourceLabels:
        - __name__
      - action: drop
        regex:
transformation_(transformation_latencies_microseconds|failures_total)
      sourceLabels:
        - __name__
      - action: drop
        regex:
(admission_quota_controller_adds|crd_autoregistration_controller_work_duration|APIServiceOpenAPIAggregationControllerQueue1_adds|AvailableConditionController_retries|crd_operations_total)
      sourceLabels:
        - __name__
      port: https-metrics
      relabelings:
        - sourceLabels:
            - __metrics_path__
          targetLabel: metrics_path
      scheme: https
      tlsConfig:
        insecureSkipVerify: true
    - bearerTokenFile: /var/run/secrets/kubernetes.io/serviceaccount/token
      honorLabels: true
      honorTimestamps: false
      interval: 30s
      metricRelabelings:
        - action: drop
```

```
regex: container_(network_tcp_usage_total|network_udp_usage_total|tasks_state|cpu_load_average_10s)
sourceLabels:
  - __name__
path: /metrics/cadvisor
port: https-metrics
relabelings:
  - sourceLabels:
    - __metrics_path__
    targetLabel: metrics_path
scheme: https
tlsConfig:
  insecureSkipVerify: true
- bearerTokenFile: /var/run/secrets/kubernetes.io/serviceaccount/token
  honorLabels: true
  interval: 30s
  path: /metrics/probes
  port: https-metrics
  relabelings:
    - sourceLabels:
      - __metrics_path__
      targetLabel: metrics_path
    scheme: https
    tlsConfig:
      insecureSkipVerify: true
jobLabel: k8s-app
namespaceSelector:
  matchNames:
    - kube-system
selector:
  matchLabels:
    k8s-app: kubelet
```

Apply, and it's done.

```
kubectl apply -f kubelet-servicemonitor.yaml
```

Traefik

I do not use Traefik much in my setup, but it is there, and it also exposes Prometheus-ready data, so why not...

Create file: `traefik-servicemonitor.yaml`:

```
apiVersion: monitoring.coreos.com/v1
kind: ServiceMonitor
metadata:
  labels:
    app: traefik
    release: prometheus
    name: traefik
  name: traefik
  namespace: monitoring
spec:
  endpoints:
    - port: metrics
  namespaceSelector:
    matchNames:
      - kube-system
  selector:
    matchLabels:
      app: traefik
```

And apply:

```
kubectl apply -f traefik-servicemonitor.yaml
```

OpenFaas

Don't worry about this one, if you deployed as described in my guide/notes you already have Prometheus set up and collecting data.

```
root@control01:~/home/ubuntu/monitoring# kubectl get pods -n openfaas
NAME                                READY   STATUS    RESTARTS   AGE
alertmanager-7b849dbf96-hkkg6       1/1     Running   0           7d18h
basic-auth-plugin-6dc5dd997f-7vq6j  1/1     Running   0           7d18h
gateway-8d768fb57-4rqnr             2/2     Running   0           35d
nats-7fd76d6d5-6677w               1/1     Running   0           35d
prometheus-54c55c4fd-89g9k         1/1     Running   0           35d
queue-worker-64b75867c8-kn7br       1/1     Running   0           35d
```

We will point Grafana to suck data from this instance later as well.

Done for now!

I know, so much text! I could probably print the whole K3s Kubernetes cluster setup as a book, drop it on somebody, and it would flatten them 🙄.

We should now have the following Service Monitors up and ready to be scraped by Prometheus.

```
root@control01:~/home/ubuntu/monitoring/kube-state-metrics# kubectl get ServiceMonitor -n monitoring
NAME                                AGE
kube-state-metrics                 14d
kubelet                           12d
longhorn-prometheus-servicemonitor 15d
node-exporter                     15d
traefik                           12d
```

Phew! we have the hardest and longest part behind us; you deserve a drink, and if you found this useful, help me to get one too. I would appreciate that a lot.

👍 Liked it ? Buy me a drink :)

Move on to [Prometheus](#)

Last update: October 20, 2021

Comments

What do you think?

2 Responses

👍

😄

😍

😮

😡

😞

Upvote

Funny

Love

Surprised

Angry

Sad

0 Comments

<https://lrpi4cluster.com>

Disqus' Privacy Policy

Login

🤍 Favorite

Tweet

Share

Sort by Best

Start the discussion...

LOG IN WITH

OR SIGN UP WITH DISQUS ?

Name