

Deployment

Let's deploy Grafana to read data from Prometheus instances.

This will be simple. In our `monitoring` directory, create a new subdirectory called `grafana`, and in it we will create following files:

grafana-pvc.yaml

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: longhorn-grafana-pvc
  namespace: monitoring
spec:
  accessModes:
    - ReadWriteOnce
  storageClassName: longhorn
  resources:
    requests:
      storage: 10Gi
```

This will be our persistent storage, it's to keep the dashboards saved. As far as I understand it, Grafana does not keep the data, so we don't have to have so much space dedicated to it (mine is using like 400MB).

grafana-deployment.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app: grafana
  name: grafana
  namespace: monitoring
spec:
  replicas: 1
  selector:
    matchLabels:
      app: grafana
  template:
    metadata:
      labels:
        app: grafana
    spec:
      containers:
        - env: []
          image: grafana/grafana:latest
          name: grafana
          ports:
            - containerPort: 3000
              name: http
          readinessProbe:
            httpGet:
              path: /api/health
              port: http
          resources:
            limits:
              cpu: 200m
              memory: 200Mi
            requests:
              cpu: 100m
              memory: 100Mi
          volumeMounts:
            - mountPath: /var/lib/grafana
              name: grafana-storage
              readOnly: false
      nodeSelector:
        node-type: worker
      securityContext:
        fsGroup: 65534
        runAsNonRoot: true
        runAsUser: 65534
      serviceAccountName: grafana
      volumes:
        - name: grafana-storage
          persistentVolumeClaim:
            claimName: longhorn-grafana-pvc
```

Fairly standard deployment, I mentioned most of the "kinks" I use before, like `nodeSelector` etc...

grafana-serviceAccount.yaml

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: grafana
  namespace: monitoring
```

Just a service account for Grafana.

grafana-service.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: grafana
  namespace: monitoring
spec:
  selector:
    app: grafana
  type: LoadBalancer
  ports:
    - name: http
      port: 3000
      targetPort: http
  loadBalancerIP: 192.168.0.236
```

Classic for us by now. I'm creating external IP for Grafana to run on 192.168.0.236, and port 3000.

Jump one folder up, and apply to the whole folder:

```
cd ..
kubectl apply -f grafana/
```

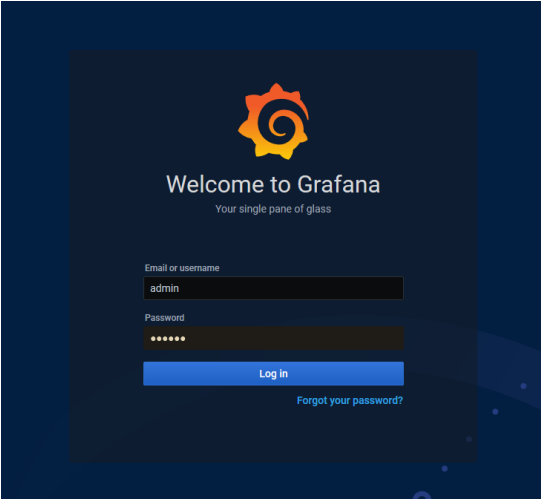
Check if grafana pod is deployed:

```
root@control01:~/home/ubuntu/grafana# kubectl get pods -n monitoring
NAME                                READY  STATUS   RESTARTS  AGE
grafana-5b799b4c8c-qzp52            1/1    Running  0          8d
.
```


Basic setup

You should be able to connect to the IP of Grafana now.


Default login and password is `admin:admin`



Then, go down and change your account name, password etc...



Next, we need to define the source where Grafana should look for data.



Click on 'Add data source' and choose 'Prometheus', a new tab with settings will pop up. Set a name for your instance, for example 'Prometheus-main'. This is so we can differentiate sources later. The next important value is 'URL'. If you remember, back when we deployed the Prometheus file 'prometheus-service-local.yaml', we created ClusterIP, and in another file, MetalLB IP. You can choose any of them. To check look at the services:

```
root@control01:~/home/ubuntu/grafana# kubectl get services -n monitoring
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
-	-	-	-	-	-
-	-	-	-	-	-
prometheus	ClusterIP	10.43.117.147	<none>	9090/TCP	14d
prometheus-external	LoadBalancer	10.43.49.187	192.168.8.235	9090:38850/TCP	14d
-	-	-	-	-	-
-	-	-	-	-	-

So, in the 'URL' either put IP or NAME, so for example, using internal ClusterIP, entering 'http://10.43.117.147:9090' should work.

At the bottom click 'Save & Test'; it should check and save the data source.

Add another data source; this will be for OpenFaaS (if you have it). Same drill as above; just check your IP for OpenFaaS Prometheus.

```
root@control01:~/home/ubuntu/grafana# kubectl get services -n openfaas
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
-	-	-	-	-	-
-	-	-	-	-	-
prometheus	ClusterIP	10.43.238.226	<none>	9090/TCP	35d

So use '10.43.238.226:9090' for 'URL', and name it 'Prometheus-OpenFaaS', or something that will let you know it's OpenFaaS data.


Note

I just noticed that both Prometheus instances have the same name for ClusterIP = 'prometheus', therefore I opt for IP instead of 'http://prometheus:9090', as I'm not sure if the internal DNS would mess something up.

Some Graphs

My final goal is to create my own dashboard with data I want. But before we get to that, we can use an already existing collection (and later pick and choose what we want from them).

Click on the plus sign and then 'Import':



Next, type '8171 info Import via grafana.com'.

Import via grafana.com

8171

Load

Import via panel json

Load

Where did I get the ID? Well, here: [Grafana Dashboard](#).

Click Load.

In the next window, name the dashboard if you like, but more importantly, choose the source for your main Prometheus instance.

Importing Dashboard from Grafana.com

Published by

kiddouk

Updated on

2018-10-01 08:52:34

Options

Name

Kubernetes Nodes

Folder

General

Unique identifier (uid)

The unique identifier (uid) of a dashboard can be used for uniquely identify a dashboard between multiple Grafana installs. The uid allows having consistent URLs for accessing dashboards so changing the title of a dashboard will not break any bookmarked links to that dashboard.

dUMN5x0mk

Change uid

Prometheus

Prometheus-monitoring

Import

Cancel

Click Import.

Tadaaaa! Your first graphs. It should take you to them immediately, and you can choose data from a specific server on the top.



Here is a list of other dashboards that work, mostly, out of the box.

[Kubernetes](#) [Kubernetes Longhorn](#) [OpenFaaS 1](#) [OpenFaaS 2](#)

And that's really it. In the next chapter, I look into creating my own dashboard, but until then I need to have something to drink and chill out.

🍷 Liked it ? Buy me a drink :)

Last update: May 26, 2022

Comments

What do you think?

4 Responses

👍

😂

❤️

😮

😡

😢

Upvote

Funny

Love

Surprised

Angry

Sad

0 Comments

<https://rpi4cluster.com>

[Disqus' Privacy Policy](#)

[Login](#)

🍷 Favorite

🐦 Tweet

📱 Share

Sort by Best

Start the discussion...

LOG IN WITH

OR SIGN UP WITH DISQUS ?

Name